Honors Draft

Raven McKnight

2/9/2020

Abstract

This will be my abstract

1 Introduction

Public transit ridership has been in decline around the country for years, but particularly since the mid 2010s (). Following a spike in ridership related to the 2008 housing market crash, ridership has been steadily trending downwards in most metropolitan areas (). Bus ridership is particularly affected compared to more "desirable" light rail and commuter rail. Transit agencies are naturally interested in understanding the specific causes of these declines, as well as more generally being able to predict transit demand throughout space and time.

Transit ridership, or demand, is often forecasted using models from the policy and planning world, such as the classic four-step model. Four-step transportation demand models consist of simulating trips and their spatial distribution before assigning them modes (ie bus versus car versus bicycle) and routes. Other models, called activity-based models predict transportation demand based on when riders are likely to conduct various tasks (commuting, shopping, recreation, etc). Others still are based largely on land-use characteristics under the general premise that different types of development (or lack thereof) are more or less likely to draw riders.

It is somewhat rarer to see more traditional statistical models applied to transit ridership. In the context of predicting future ridership, the above methods are often sufficient. Additionally, they are the industry standard and allow for the integration of expert planning knowledge. However, in the context of understanding *past* ridership, we can apply more explanatory methods. Specifically, we can use hierarchical spatial Bayesian models to incorporate many covariates in addition to the spatial structure that necessarily underpins transit ridership (ie, boardings can only occur at existing bus stops).

Hierarchical spatial models allow us to understand the patterns of ridership within a city with great granularity. Some existing literature explores the decline of transit ridership *between* cities. Few studies have been conducted on the various patterns of ridership *within* a single metropolitan area. Additionally, few studies incorporate many demographic predictors. Naturally, we cannot collect demographic information for every rider of a transit system. Instead, we can use demographic predictors associated with spatial units such as census block groups to explore the question of *who* rides – or, at the very least, *where* do they ride.

2 Background

2.1 Metro Transit

Metro Transit is the primary transit provider in the Minneapolis-Saint Paul metropolitan area. The agency operates one commuter rail line (Northstar), two light rail lines (Blue Line and Green Line), and over 100 bus routes. Since the 2014 opening of the Green Line, rail ridership has increased each year. Bus ridership, however, has been in decline. According to the Metro Transit Riders' Almanac blog, much of the decline is bus ridership can be attributed to the busiest urban local routes as well as situational factors, such as lower-than-average gas prices and a system-wide fare increase in 2017.

The agency is naturally interested in understanding with more detail the reasons for and nature of these declines.

Conducted in tandem with Metro Transit's Network Next bus system redesign, this study...

2.2 Transit Market Areas

Metro is interested in predicting transit demand across the region in order to guide route planning and service provision. Currently, Metro Transit uses a set of 5 Transit Market Areas to estimate spatial demand. The Transit Market Areas are writen into the official Transportation Planning Policy (TPP) and are calculated using a simple linear regression. Following the notation of the TPP, the formula for determing Transit Market Areas is expressed

$$\begin{aligned}\text{Transit Market Index } &= 0.64 * \text{Population Density} + 0.20 * \text{Employment Density} + \\ &\quad 0.23 * \text{Intersection Density} + 0.11 * \text{Automobile Availability}\end{aligned}$$

where each predictor is logged and scaled by developed land acreage per census block group. Here, automobile availability refers to the number of adults over age 16 less the total number of automobiles available in a block group (scaled by developed land acreage). There is an additional indicator variable, omitted in this notation, for the census block group containing the MSP International Airport. For more documentation of the official Transit Market Areas, refere to the TPPs Appendix G.

Census block groups are split into 5 market areas based on the Transit Market Index described above, where the highest market area (Market Area 1) is expected to support high-frequency, all-day service and the lowest market area (Market Area 5) is expected to support peak commuter express service and park-and-rides, if that.

The Transit Market Index values are geographically smoothed to create more-or-less concentric Transit Market Areas. The linear regression is an intuitive way to think about transit ridership across space. The four predictors used are common-sense indicators of transit ridership: high population and employment density are characteristic of trip origins and destinations, automobile availability is reasonably assumed to be related to transit ridership, and intersection density is a proxy for the "walkability" of an area. However, the exsiting Transit Market Area model may be over-simplifying the complex question of transit ridership.

3 Data

3.1 Metro Transit Data

This analysis relies on several Metro Transit provided data sets. The two primary sources are automatically reported by in-service vehicles, yielding billions of rows of observations.

Automatic Passenger Count (APC) data is reported every time a bus opens its doors. Two beams of light detect movement through the doors of the bus and counts the number of passengers getting on and off at each bus stop. Naturally, this data source is flawed: sometimes, vehicles fail to report data, and the sensors can be easily tricked. Someone getting off the bus pulling a suitcase behind them, for example, would likely be counted as two passengers exiting. However, APC is the most granular ridership data available, giving us the most control over our spatial aggregations of ridership. This is the primary data source used in this analysis.

Automatic Vehicle Location (AVL) data is similarly reported by in-service vehicles (approximately) every 8 seconds. The vehicle reports in GPS location continuously while in service. We use AVL data to adjust for missing APC data.

Metro Transit's Schedule as *planned*, rather than as *run*, is also used to correct missing APC data.

This analysis began by pulling all APC, AVL, and schedule data from 2015-2018. The initial data pull consisted of approximately 496 *billion* rows of data. The analysis presented here focuses on 2017 data.

3.1.1 Data Interpolation

In theory, we have APC data for each run bus trip in 2017. However, we know that, for a variety of reasons, this is untrue. The automatic data reporting technology on Metro Transit buses is flawed and can fail to

report data unpredictably. Additionally, the trips missing data may not be random: all buses of one "series", for example, may fail to report at once. This could lead to the APC dataset understimating boardings on a particular set our routes or locations.Therefore, we use the three data sets described above to create a more complete augmented APC data set.

The algorithm used to interpolate data works essentially as follows:

1. Compares scheduled trips to trips we have APC data for. If a trip was scheduled but has no APC records, we consider that trip missing.
2. For missing trips, check AVL data. If there are no AVL records, we consider the trip cut (ie, we assume the trip was not run).
3. For missing trips with AVL data, we estimate boardings for that trip by taking the average number of boardings for trips of that nature (trips on the same route, day of week, season, location, etc).

3.1.2 Data Aggregation

For the models presented in this study, we aggregate boardings at individual bus stops to the census block group. This is in large part to match the granularity of the American Community Survey covariates we are interested in using.

Additionally, we aggregate boardings counts to average weekday boardings. This a) greatly reduces the number of rows of data in play, and therefore shortens computation time significantly, and b) allows us to build a spatial model without simultaneously incorporating temporal trends. The data interpolation described above will be more impactful in future studies incorporating temporal trends.

3.2 Other data sources

All other data sources used in this study are from outside sources. Census Block Groups are used as the areal unit of analysis. Covariates of interest are primarily from the 2017 American Community Survey (ACS) 5-Year Estimates. Additional variables related to employment are sourced from the Longitudinal Origin-Destination Employment Statistics (LODES) dataset. A full list of covariates used in this study is below.

4 Models

4.1 Poisson Regression

For each block group $i = 1, 2, 3, ..., 1495$ with any ridership, the "baseline" model is a simple Poisson regression. Poisson regressions are a type of generlized linear model with the natural log as its link function. For ridership in Y_i and covariates $x_1, ..., x_k$ discussed above, the Poisson regression is written

$$\begin{aligned} Y_i &\sim \text{Poisson}(E_i\lambda_i) \\ \eta_i = log(\lambda_i) &= \beta_0 + \sum_{k=1}^{K} x_k^T \beta_k \\ \beta_0, \beta &\sim \text{Normal}(0, 1) \end{aligned}$$

where β_0 is the intercept and E_i is an offset term.

The offset E_i is often termed "exposure" in applications such as disease risk modeling where the offset may be the number of observed cases in a previous year, etc. More technically, the offset term scales the Poisson output to be a rate rather than a count. This is appropriate when observations i have different potentials for response Y. For example, a county with higher population will naturally have a higher count of patients with asthma than a county with a smaller population. Mathematically, the offset is a covariate with parameter set equal to 1. Recall, the Poisson regression assumes we can model the mean of response Y, $\bar{Y}$ with a combination of linear predictors:

$$log(\bar{Y}) = \beta_0 + \sum_{k=1}^{K} x_k^T \beta_k$$

Therefore, if we wish to model the rate Y/E, the equation is rewritten

$$log(\bar{Y}/E) = \beta_0^{'} + \sum_{k=1}^{K} x_k^T \beta_k^{'}$$

$$log(\bar{Y}) = log(E) + \beta_0^{'} + \sum_{k=1}^{K} x_k^T \beta_k^{'}$$

In the basic Poisson regression, we often give β_0 and β vague priors such as Normal$(0, 1)$. The exposure term E_i comes from the data and does not recieve a prior. In our case, we define E_i to equal the number of times a bus stops in block group i. This allows us to control for block groups with more or less supply.

4.2 Overdispersed Poisson Regression

The most obvious assumption made in a Poisson regression is that response variable Y follows a Poisson distribution. In addition to requiring integer counts, this assumption requires that $\text{Var}(Y) = \text{E}(Y)$. This is often not true in practice. In the case of the Metro Transit ridership data, $\text{Var}(Y) \approx 7 * \text{E}(Y)$. In this case, we call the unexpectedly large variance "overdispersion" or "extra-Poisson variance."

Data with this feature can be difficult to model directly with a Poisson regression. One option is to use a generalization of the Poisson distribution, such as a Negative-Binomial distribution, as the observation model. For the Metro Transit case study, we instead modify the latent function to include a seperate term to account for overdispersion. As such, the Poisson regression above can be rewritten to include

$$\eta_i = log(\lambda_i) = \beta_0 + \sum_{k=1}^{K} x_k^T \beta_k + \theta\sigma$$
$$\theta \sim \text{Normal}(0, 1)$$
$$\sigma \sim \text{Normal}(0, 5)$$

where θ is a set of heterogenous random effects. In practice, θ is scaled by its variance parameter σ to allow eassier fitting in Stan. The addition of this set of random effects improves fits on data with overdispersion. With sufficiently vague priors, θ and σ can account for most or all of the overdispersion not accounted for by $\beta_0 + \sum_{k=1}^{K} x_k^T \beta_k$.

4.3 Horseshoe Priors and Variable Selection

In the baseline model, we give coefficients $\beta_1, ..., \beta_k$ simple normal priors. With a large set of possible covariates, however, we may reasonably expect some of the β_k coefficients to be equal to zero – in other words, that some of the parameters x_k have no effect on response Y. In frequentist applications, we might use a method such as the LASSO (Least Absolute Shrinkage and Selection Operator) to identify relevant variables. There are several Bayesian alternatives for variable selection with generally consist of applying particular priors to β.

4.3.1 Spike and Slab - needs graphics

In the Bayesian setting, there are two primary methods for variable selection. The first is called the "spike-and-slab" prior (Mitchell and Beauchamp, 1988; George and McCulloch, 1999) and has often been considered the "gold standard" for sparse Bayesian regression, or variable selection (Piironen and Vehtari, 2017). This prior is often expressed as

$$\beta_j|\lambda_j, c, \epsilon \sim \lambda_j \text{Normal}(0, c^2) + (1 - \lambda_j)\text{Normal}(0, \epsilon^2)$$
$$\lambda_j \sim \text{Bernoulli}(\pi)$$

for $j = 1, 2, ..., J$ where $\lambda_j \in 0, 1$ indicates whether β_j is from the "spike" (ie, near zero) or the "slab" (ie, nonzero). The "spike" is the area where most of the prior density for β is centered (generally around 0) while the "slab" refers to the low-density extent of the prior. The spike-and-slab prior encourages all β_j towards zero; only β_j sufficiently far from 0 will be estimated to be from the slab.

The primary challenge with the spike-and-slab (and other Bayesian shrinkage methods) is the selction of priors. Priors must be set for the width of the slab c and the prior inclusion probability π. Here, π reflects our prior understanding of the sparsity of β. In practice, we rarely have strong prior knowledge of the number of predictors we expect to be distinguishable from zero. This can make the implementation of sparse regression more challenging than, say, the frequentist LASSO.

4.3.2 Horseshoe - needs graphics

The second method for Bayesian sparse regression is to give β some continuous *sparsity inducing prior*. **Van Erp et al (2019)** provide an introduction to many such continous priors. Ridge regression and the Bayesian LASSO are two of the most approachable priors, but all sparsity inducing priors utilize the same logic as the spike-and-slab: give β a prior with *most* of its mass near 0, to shrink irrelevant coefficients, and the rest of its mass "far" from 0. Identifying "far" depends more or less on the specific data and model in question.

The horseshoe prior, proposed by **Carvalho et al, 2010** is a popular choice in Bayesian literature, in part because it is similar to the "gold standard" spike-and-slab. Following the notation of **Vehtari & Piironen**, the horseshoe prior can be expressed

$$\beta_k \mid \lambda_k, \tau \sim \text{Normal}(0, \tau^2\lambda_k^2)$$
$$\lambda_k \sim \text{Half-Cauchy}(0, 1)$$

for $k = 1, 2, ..., K$. The horseshoe prior is so-named because for fixed values $\tau = \lambda_k = 1$, the prior resemebles a Beta(1/2, 1/2) distribution, or a horseshoe.

4.3.3 Regularized Horseshoe - needs graphics

The horseshoe is often favored because it is a global-local shrinkage prior. This means that τ shrinks *all* parameters towards 0 while the local parameter λ_k and the heavy Cauchy tails allow larger coefficients to remain unshrunk. While this is precisely the goal of a sparsity inducing prior, the horseshoe prior can fail to regularize large coefficients *at all*, which can be a problem when parameters are weakly identified by data. With no regularization applied to the largest coefficients, there is a risk of overfitting. Additionally, there is no consensus in the literature for assigning priors to *tau*. Piironen and Vehtari (2017) introduced the *regularized horseshoe* to address both of these shortcomings of the horseshoe. The regularized horseshoe builds upon **need to figure out equation numbering** such that

$$\beta_k \mid \lambda_j, \tau, c \sim \text{Normal}(0, \tau^2\tilde{\lambda}_k^2)$$
$$\tilde{\lambda}_k^2 = \frac{c^2\lambda_k^2}{c^2 + \tau^2\lambda_k^2}$$
$$\lambda_k \sim \text{Half-Cauchy}(0, 1)$$
$$c \sim \text{Inverse-Gamma}(a, b)$$

where $c > 0$ helps to regularize β_k far from zero. When c approaches infinity, the regularized horseshoe becomes the standard horseshoe. Additionally, c in the regularized horseshoe corresponds to the slab width in the spike-and-slab prior.

As with the spike-and-slab, the selection of priors is a challenge with the regularized horseshoe. Piironen and Vehtari (2017) provide a set of recommendations for setting priors which we follow in this study.

likely need to say more here

4.4 Adding Spatial Structure

When modeling data with a spatial component, we generally expect adjacent areas to be more similar than areas which are far apart. While this is a straightforward assumption, it can improve model fits in several ways. First, it can encode prior information about response Y not otherwise represented by covariates x (ie, that nearby observations are more similar than far-flung ones). Second, it can provide geographic smoothing when observations are sparse or noisy. This is often the case in small areas or when observing events which can only occur at specific locations, such as air quality measured by sensors.

In the case of Metro Transit's ridership data, we expect encoding spatial structure to improve the model fit in several ways. First, we know that assigning transit ridership to census geography is inherently flawed. This is because census geographies, including block groups, are generally bounded by major roads (in addition to natural features such as lakes and political boundaries such as city limits). Naturally, bus stops tend to exist along those same major roadways.

This bounding convention means that boardings are often assigned to different block groups depending on which direction the passenger is travelling. For example, consider the A-Line running along Snelling Avenue. Boarding at Snelling and Grand *northbound* puts you in block group _____ while boarding at the same intersection but *southbound* puts you in block group _____. Integrating spatial structure will help to smooth ridership across block group bounds.

Additionally, we know intuitively that transit ridership does not occur in a single block group. Particularly in dense areas, it's very likely for riders to board in a block group other than the block group they reside in. For example, I live in block group _____ but walk about 10 minutes to board the A-Line southbound at Snelling and Randolph. Commuting patterns such as these can assign ridership to different block groups than we might expect. We hope that the geographic smoothing provided by spatial modeling will allow our estimates or ridership to better represent our understanding of ridership. ***oof***

4.4.1 Conditional Autoregressive Priors

Conditional Autoregressive (CAR) priors are one of the most widespread Bfayesian methods for modeling spatial autocorrelation. CAR models were introduced in Besag 1974 and remain perhaps the primary method for Bayesian areal data modeling. Note that CAR models are recommended for use as *priors* rather than as an observation model.

Areal data corresponds to finitely many discrete areal units, such as counties or census tracts. CAR priors are often used in the context of small-area count data. This sort of response variable tends to be noisy, particularly when the counted even is rare, the population of areal unit i is small, or the physical boundaries of areal units present challenges. CAR priors were designed in part to smooth such noisy counts. As such, they are often applied in epidemiological disease risk modeling in which disease occurences may be rare.

Generally, CAR models rely on a binary neighborhood structure. Areal units i and j are considered neighbors if they share a boundary. For strictly rectangular lattice data, we often choose either a "Queens" or "Rooks" neighborood structure. For irregularly shaped areal data, such as our 1495 census block groups, we general default to a queen neighborhood structure wherin areal units that share *any* points of contact are deemed neighbors.

For n regions, the neighorhood relationships are encoded in an $n \times n$ neighborhood matrix W. Matrix W is defined such that $w_{ij} = 1$ if if regions i and j are neighbors (denoted $i \sim j$) and 0 otherwise. This matrix is symmetric. Note that regions are not considered neighbors of themselves, ie the diagonal entries $w_{ii} = 0$.

The spatial interactions are modeled by a random variable ϕ. Here, ϕ is a vector $\phi = \phi_1, \phi_2, ..., \phi_n$. The distribution of each ϕ is determined by the sum of its neighbors' values such that

$$\phi_i \mid \phi_j, j \neq i \sim \text{Normal}(\sum_{i=1}^{N} w_{ij}\phi_j, \sigma^2)$$

The conditional distribution of each ϕ_i is helpful in building intuition about the functionality of the CAR prior. Basing the value of each ϕ_i off of its neighbors is how the CAR priors performs spatial smoothing.

Besag (1974) proved that the joint distribution of ϕ is a multivariate normal centered at 0 where its variance is given by a symmetric positive definite precision matrix Q. Matrix Q is simply the inverse of the covariance matrix $\sum$. This finding simplifies the above to

$$\phi \sim \text{Normal}(0, Q^{-1})$$

Precision matrix Q can be defined for multivariate normal ϕ using two other NxN matrices: the neighborhood matrix W as well as diagonal matrix D in which d_{ii} indicates the number of neighbors region i has. All off-diagonal entries are 0. Using these matrices, Q is defined

$$Q = D(I - \alpha W)$$

where I is the identity matrix and $\alpha \in (0, 1)$ determines the amount of spatial autocorrelation present in Y. Parameter α is a key component of CAR models. At $\alpha = 0$, the model assumes spatial independence and at $\alpha = 1$, perfect spatial autocorrelation.

Following Morris et al, the log probability density of ϕ is proportional to

$$\frac{M}{2}\log(\det(Q)) - \frac{1}{2}\phi^T Q\phi$$

Calculating this value is computationally expensive. For models with 1,000 areal units, for example, calculating the determinant requires 1 billion operations (Morris et al, 2019). Given that the Metro Transit ridership data corresponds to 1,495 areal units, a strict CAR prior may be too computationally inefficient for our models.

4.4.2 Intrinsic Conditional Autoregressive

The Intrinsic Conditional Autoregressive (ICAR) prior is a slight simplifation of a CAR prior. ICAR priors set $\alpha = 1$ which simplifies the definition of Q

$$Q = D - A$$

thereby setting $det(Q) = 0$. As a result, the first term in **equation numbering** can be simplified to

$$-\frac{1}{2}\phi^T Q\phi$$

This reduces the computational expense of computing CAR models significantly. Notably, the ICAR prior is improper but yields proper posteriors (Morris et al, 2019).

The ICAR prior specifies each ϕ_i to be normally distributed with its mean equal to the mean of its neighbor's values. This is how the CAR/ICAR model "borrows strength" from geographic neighbors to smooth noisy estimates. Intuitively, the variance of each ϕ_i decreases as its number of neighbors d_i increases. The conditional distribution for each ϕ_i can be written

$$\text{p}(\phi_i \mid \phi_{i\sim j}) = \text{Normal}(\frac{\sum_{i\sim j}\phi_i}{d_i}, \frac{\sigma_i^2}{d_i})$$

Here, the variance σ is unknown. The conditional specification of each ϕ_i is helpful in building intuition about the assumptions the ICAR component makes and how it performs its "smoothing". The ICAR prior matches our intuition about spatial autocorrelation: areal unit i is similar to the areal units surrounding it, and the more areal units surround it, the more confident we feel in that similarity. The joint distribution can be rewritten

$$\mathrm{p}(\phi) \propto \exp(-\frac{1}{2}\sum_{i\sim j}(\phi_i - \phi_j)^2)$$

This pairwise difference specification is also helpful in building intuition about how ICAR models perform smoothing. Each $(\phi_i - \phi_j)^2$ works as a penalty term: the ???. Additionally, the pairwise difference specification is more computationally efficient to fit in Stan (Mitzi).

4.4.3 Besag-York-Mollie (BYM)

The specific ICAR model used in this analysis is the Besag-York-Mollie (BYM) Poisson model. The BYM is a "classical" spatial Bayesian method (Riebler et al, 2016). It is generally used in disease mapping studies to model rare events (ie, disease occurrence) in small areal units. The BYM capitalizes on the CAR prior's ability to smooth noisy estimates and share information across geographic units.

The latent function in the BYM model contains heterogeneous random effects as well as a spatial ICAR term in order to account for both spatial and non-spatial heterogeneity. Note that in Section 1.1, we introduced a random effect term θ. Therefore, the BYM model can be though of as decomposing θ into spatial and non-spatial components.

Specifically, the BYM model replaces η_i from equation *1* with

$$\eta_i = \beta_0 + \sum_{k=1}^{K} x_k^T \beta_k + \theta + \phi$$

where ϕ is the addition: an ICAR component. Decomposing overdispersion as such is helpful in furthering our understanding of response Y. When we use a simple random effect θ alone, we can accomplish a good model fit. However, θ provides no additional information about Y or why it varies so much. The decomposition of θ in the BYM allows us to better quantify how much variance is white noise (ordinary random effects) versus some unmeasured confounding correlated across space.

The BYM model as written above is appealing in its simplicity. However, the lack of informative hyperpriors specified for θ and ϕ can make the model incredibly challenging to fit. In theory, extra-Poisson variance could be explained 100% by θ, 100% by ϕ, or anywhere in between. In *practice*, we likely have limited information about where that ratio falls. Riebler et al (2016) further explain the sampling issues faced by the BYM model faced with no hyperpriors. In short, the sampler is forced to explore all possible combinations of θ and ϕ, no matter how unlikely a particular combination my be. In the context of fitting models with Stan, this is liekly to cause uneccesarily long computation times and perhaps to yield biased posterior estimates.

There are some suggestions in the literature for setting hyperpriors for θ and ϕ (Besag and Mollie, 1991; Clayton and Montomoli, 1995). The existing methods, however, tend to rely on the specific data at hand which makes the selection of hyperparameters unnecessarily time consuming.

4.4.4 BYM2

The BYM2 model proposed by Simpson et al and further described in Riebler et al (2016) aims to solve these sampling problems. The BYM2 takes a more "fully Bayesian" approach to hyperprior/hyperparameter selection. The primary difference between the BYM and BYM2 models is the addition of a mixing parameter, ρ. The BYM2 model rewrites the BYM as

$$\eta_i = \beta_0 + \sum_{k=1}^{K} x_k^T \beta_k + ((\sqrt{\frac{\rho}{s}})\theta^* + ((\sqrt{1-\rho})\phi^*)\sigma$$

where $\rho \in (0, 1)$ determines how much variance/overdispersion is caused by spatial versus non-spatial error terms. Like the popular Leroux CAR prior, proposed by Leroux et al (2009), the BYM2 scales both θ and ϕ by σ, the standard deviation of the combined error terms (Morris et al, 2019). In this parameterization, s is a scaling factor such that $\text{Var}(\theta_i) \approx \text{Var}(\phi_i) \approx 1$. The equal unit variance is necessary for σ to truly be the standard deviation of the error terms.

Setting priors is somewhat more straightforward for the BYM2. The ICAR component, ϕ^* remains unchanged. Riebler and Morris recommend the prior $\theta \sim \text{Normal}(0, n)$ where n is the number of connected graphs in the neighborhood graph. In many cases, such as modeling data for all counties in a state, $n = 1$. When $n > 1$, the variance is different in each subgraph which affects σ and s. It is possible to fit the model in this case, although computation time is much longer.

A relatively vague Normal(0, 1) prior is appropriate for σ. Riebler et al (2016) also propose either a Beta(1/2, 1/2) or more specialized, complex prior for mixing parameter ρ. For this study, we use the simpler Beta(1/2, 1/2) prior.

5 Results

In this study, we fit several increasingly complex models building upon the baseline Poisson regression, Model 1. Model 2 incorporates a simple set of random effects to account for overdispersion. Model 3 incorporates regularized horseshoe priors for variable selection. Model 4 includes the spatial ICAR priors but no horseshoe priors, and Model 5 includes both regularized horseshoe priors and ICAR priors.

All five models were fit using the programming language Stan. Stan is a probabilistic language which uses Hamiltonian Monte Carlo (HMC) and a specialized No U-Turn Sampler (NUTS) to compute joint log probability densities. The specifics of Stan are beyond the scope of this paper. However, a few of its particular diagnostics are helpful in understanding the differences between the five model fits.

6 Discussion

- bymhrs is probably not "worth it", needs to be reparameterized/coded by someone better than me
- incorporating overdispersion is the most important improvement in terms of model fit